

1. Dataset

The Facebook Egonets dataset represents friendship networks collected from real Facebook users. Each network is an ego network, meaning it contains:

- An ego user
- Their friends (nodes)
- The friendship links between those friends (edges)

Dataset summary:

- **4,039 nodes (users)**
- **88,234 edges (friendships)**
- No real bots → we inject **271 synthetic bots** based on structural anomalies:
 - **Degree above average + very low clustering**
Nodes with many connections but almost no triangles, behaving unlike real social users.
 - **Low eigenvector centrality**
Nodes connected to weak or unimportant parts of the network.
 - **Low betweenness centrality**
Nodes that rarely lie on important shortest paths, meaning they do not bridge communities.
 - **Low closeness centrality**
Nodes that are far from the rest of the network, structurally isolated.

Nodes meeting these anomaly conditions were labeled as synthetic bots → 271 total bots.

2. Graph Metrics

For each node we compute:

a) Degree

- Measures the number of direct connections.

b) Clustering coefficient

- Shows how strongly a node's neighbors are interconnected.

c) Centrality measures

- **Eigenvector centrality:** influence based on neighbors' influence

- **Betweenness centrality:** how often a node lies on shortest paths
- **Closeness centrality:** how close a node is to all others

d) Community detection

- We use the Louvain algorithm to assign each node to a community.

```
PS C:\Users\MALAK\Desktop\Social networks-Assignment2> python test.py
  node  degree  clustering  eigenvector  betweenness  closeness  community
  0      0    347  0.041962  3.391796e-05  0.169475  0.353343      4
  1      1     17  0.419118  6.045346e-07  0.000000  0.261376      4
  2      2     10  0.888889  2.233461e-07  0.000000  0.261258      4
  3      3     17  0.632353  6.635648e-07  0.000006  0.261376      4
  4      4     10  0.866667  2.236416e-07  0.000000  0.261258      4
```

These metrics form the core features for the classifier.

3. Interpretation of three models

Accuracy stays very high (93–96%) in all 3 scenarios because:

- The dataset is extremely imbalanced
 - 1136 humans
 - ~76 bots
- That means bots represent only ~6% of the dataset.

So even if the model completely fails to detect bots, accuracy stays high:

- If the model predicts “human” for everything →
 $\text{Accuracy} \approx 1136 / 1212 = 94\%$

So, Accuracy hides model failures as it is misleading in imbalanced classification and should not be used alone.

3.1 Baseline Bot Detection Model

We use:

- Node2Vec to generate 64-dimensional graph embeddings
- GraphSAGE (2-layer GNN) for classification
- Train/test split: 70% / 30%
- Labels: bot = 1, human = 0

The baseline model is trained on the clean graph (no attack).

```

===== BASELINE GNN RESULTS =====
Accuracy: 0.9636963696369637
      precision    recall   f1-score   support
0         0.97     0.99     0.98     1136
1         0.80     0.57     0.66      76
accuracy          0.96     0.96     0.96     1212
macro avg       0.88     0.78     0.82     1212
weighted avg    0.96     0.96     0.96     1212

```

Interpretation

- Strong performance on humans.
- Bot recall = 0.57, meaning the model detects 57% of bots and misses 43%.

3.2. Structural Evasion Attack

Objective:

- Make bots look structurally similar to humans.

Method:

- Each bot connects to high-degree influencer nodes
- Adds 7 new edges per bot
- Bots move deeper inside real communities
- Their embeddings become more "human-like"

Effect:

- Graph structure changes, but labels remain unchanged.

```

===== EVASION ATTACK RESULTS =====
Accuracy: 0.943069306930693
      precision    recall   f1-score   support
0         0.94     1.00     0.97     1136
1         0.77     0.13     0.22      76
accuracy          0.94     0.94     0.94     1212
macro avg       0.86     0.56     0.60     1212
weighted avg    0.93     0.94     0.92     1212

```

Interpretation

- Bots successfully hide inside densely connected human clusters.
- Bot recall drops from 0.57 → 0.13 (a ~77% decrease).
- The model still performs well on humans, but becomes nearly blind to bots.
- Evasion is effective, but not as damaging as poisoning.

3.3 Graph Poisoning Attack (Label Flipping)

Objective:

- Corrupt training labels so the model learns wrong rules.

Method:

- Select bots inside the training set only
- Flip 30% of bot labels from 1 → 0
- No structural changes to the graph

Effect:

- The model is trained to believe some bots are humans, destroying decision boundaries.

===== POISONING ATTACK RESULTS =====					
	precision	recall	f1-score	support	
0	0.94	1.00	0.97	1133	
1	1.00	0.03	0.05	79	
accuracy			0.94	1212	
macro avg	0.97	0.51	0.51	1212	
weighted avg	0.94	0.94	0.91	1212	

Interpretation

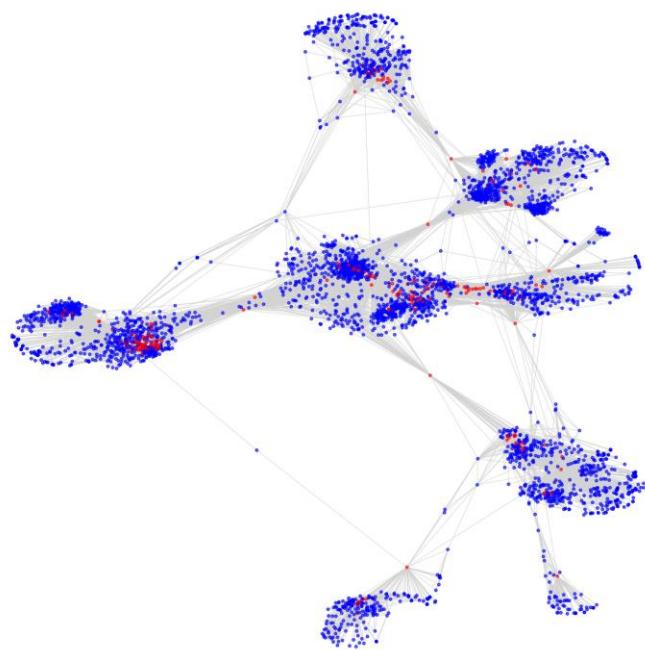
- Bot recall collapses from 0.57 → 0.03 (a 94% reduction). This means the model only detects 3 bots out of 100 on average.
- The model almost completely loses the ability to detect bots.
- Accuracy remains high due to imbalance, but this is misleading.
- Poisoning is the most destructive attack by far.

4. Visualization of Attack Scenarios:

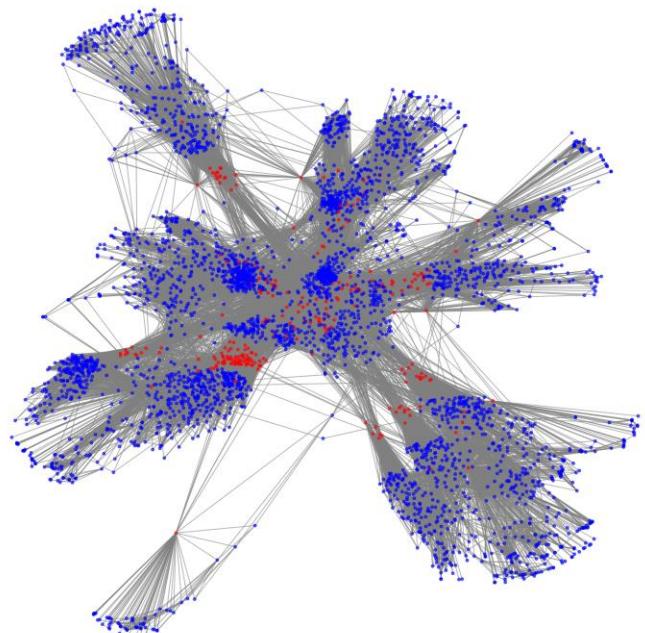
Color scheme:

- Blue = humans
- Red = bots

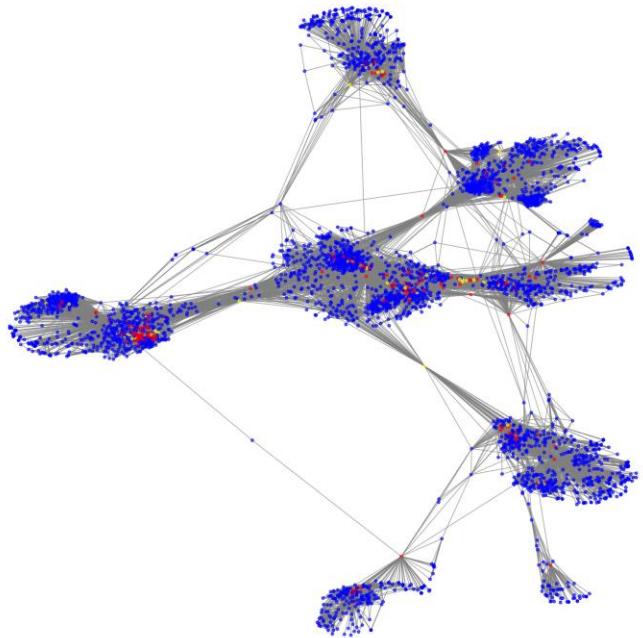
a) Baseline graph:



b) Evasion attack graph:



c) Poisoning attack graph:



- Yellow nodes = flipped nodes ($1 \rightarrow 0$)